



ADYDADDY

AI-Proof Marketer

How to save your performance-marketing job from AI, by using AI better than anyone else.

PRO

Rs 1,499 | one-time payment | instant PDF delivery

By Vibhor Sharma, Founder of AdyDaddy

A practical playbook for performance marketers.

What you get in Pro

- Everything in Starter.
- Full prompt library across research, creative, buying, and reporting.
- A weekly reporting template you can paste exports into.
- A competitor teardown checklist.

How to read this playbook

Read each chapter in order once. Then keep it open as a reference and run the prompts as you work. The value is in using it, not just reading it.

01 | MODULE

The Reality Check

What AI actually does to performance-marketing roles, and why the people who survive are the ones who learn to direct it.

Why the fear is real, but the panic is wrong

Every day there is a new headline about AI writing ads, building audiences, or 'running your entire funnel.' If your job is mostly copy-paste, upload, and report, that fear is justified. If your job is deciding what to test, why it won, and where to spend the next rupee, AI is not your replacement. It is your leverage.

What AI is genuinely good at

- Producing first drafts fast: ad copy, hooks, landing headlines, email subject lines.
- Summarising long data: turning 40 rows of ad metrics into a plain-English story.
- Generating variations: 20 hook options from one angle in seconds.
- Structuring work: account naming, test plans, briefs, checklists.
- Explaining concepts: how a bid strategy works, what a metric means.

What AI is still bad at

- Judgement under ambiguity: knowing when a 'winning' ad is a fluke.
- Accountability: owning a revenue number and living with the result.
- Context: knowing your brand, your founder, your customer, your constraints.
- Taste and brand safety: what fits the brand voice and what would embarrass the client.

TIP

Your job is moving up the stack: from doing the work, to deciding what work matters and directing AI to do the rest.

The core insight of this course

AI will not replace the performance marketer. A performance marketer who uses AI well will replace the one

who does not. This course turns you into the first kind. The job you are saving is your own, and the tool you use to save it is the same tool everyone is scared of.

02 | MODULE

AI Literacy for Marketers

How these models work well enough to use them well: strengths, failure modes, and how to stop getting generic output.

How a large language model actually thinks

An LLM is a next-word predictor trained on a huge slice of the internet. It does not know your account, your brand, or your numbers. It predicts what a confident expert would say next. That is why it sounds authoritative and why it is wrong about 15 to 30 percent of specific facts. Treat it as a brilliant, fast, sometimes-wrong junior.

The three failure modes to watch for

- Hallucination: making up a metric, a benchmark, or a 'best practice' that never existed.
- Generic output: safe, average advice because you gave it no context to be specific.
- Overconfidence: answering a question it should have pushed back on instead.

The fix for all three: context, role, and format

Generic output is almost always a prompt problem, not a model problem. The single highest-leverage habit in this course is writing prompts that give the model a role, your context, and the exact format you want back. Chapter 3 is entirely about this.

Pick one model and go deep

ChatGPT, Claude, and Gemini are all good enough. Do not chase every new release. Pick one you can pay for (the free tiers rate-limit exactly when you need them), learn its voice and its quirks, and build everything in this course on it. Switching tools weekly is a productivity tax, not a skill.

TIP

Paid tiers matter: they unlock longer context (paste a whole campaign export), better reasoning, and image and web capabilities you will use every day.

03 | MODULE

The Marketer's Prompt Playbook

The frameworks that turn AI from a novelty into a daily workhorse: role, context, task, format, and the iteration loop.

The R.C.T.F. framework

- Role: 'You are a senior performance marketer at a D2C fashion brand.'
- Context: paste the brand, the audience, the goal, the numbers you have.
- Task: one clear ask, not five bundled questions.
- Format: table, 5 bullets, a 3-line summary, a JSON list. Tell it the shape.

TIP

The more specific the role and context, the less generic the output. A one-line prompt gets a one-line answer you could have Googled.

The iteration loop: generate, judge, refine

Never accept a first draft as final. The workflow that actually ships work is: generate, read it critically, tell the model exactly what is wrong ('tone is too corporate, make it sharper, shorter, more direct'), and regenerate. Two or three rounds beats one long perfect prompt.

Prompting patterns that pay off in marketing

- 'Act as' a buyer persona, then ask what would make them click.
- 'Show me 10 versions, then rank them' to force range and comparison.
- 'Explain the reasoning first, then give the answer' for decisions.
- 'What am I missing?' after any plan, to surface blind spots.
- 'Rewrite this for X' where X is a specific audience, channel, or character limit.

Your first three reusable prompts

1. A campaign-summary prompt: paste raw ad data and get a plain-English read of what is working. 2. A

hook-machine prompt: one angle in, twenty hooks out. 3. A briefing prompt: turn a messy client note into a clean creative brief. Build these three, save them, and you already use AI daily.

04 | MODULE

AI for Research and Strategy

Market, competitor, and audience research in minutes instead of days, and turning that research into a strategy you can defend.

Competitor tear-downs without the grind

Feed the model a competitor's landing page, ad copy you have collected, and reviews, then ask for their angle, their audience assumptions, their offer structure, and the gaps they leave open. Ask it to cite what it saw so you can verify instead of trusting blindly.

Audience and messaging research

- Paste real customer reviews and ask for the language customers actually use.
- Ask for pain points ranked by how often they appear.
- Generate buyer personas, then challenge each persona with 'what is wrong with this?'
- Turn objections into ad angles and FAQ answers.

Turning research into a strategy

Have the model produce a one-page strategy: the offer, the core angle, the proof, the channel plan, and the risk. Then attack it. Ask 'what assumption here is weakest?' and 'what would make this fail?' A strategy you have stress-tested with AI is sharper than one you wrote alone.

TIP

Use the web-search mode on paid tiers for live data, but always read the sources it cites. Never ship a number you have not checked.

05 | MODULE

AI for Creative and Copy

The highest-ROI use of AI in performance marketing: producing hooks, angles, and variations at the volume testing actually needs.

The hook machine

A winning ad is usually a winning hook. Give the model your angle, your audience, and your product, and ask for 20 hooks across different styles: curiosity, pain, proof, contrarian, question, story. Then ask it to rank them by likely stop-power and explain the ranking.

From one angle to a full variation set

- Primary text in three lengths: short, medium, long-form.
- Headlines for every platform and placement.
- Ad copy mapped to each funnel stage: cold, warm, retargeting.
- Email and SMS sequences from the same core angle.

Keeping it on-brand, not generic

Paste your best-performing past copy and your brand voice notes into the prompt as reference. Ask the model to write in that exact voice. This is the difference between AI copy that sounds like everyone and AI copy that sounds like your brand.

TIP

You are still the editor. AI generates the volume; your taste picks the winners. That split of labour is exactly why you stay employed.

06 | MODULE

AI for Media Buying

Use AI for the parts of buying that are pure analysis and structure, while you keep the judgement that makes the money.

Where AI helps a media buyer

- Account structure and naming conventions that scale cleanly.
- Turning a messy export into a clean story of what is working and what to kill.
- Drafting a testing plan with a hypothesis for every variant.
- Explaining a bid-strategy or attribution concept before a client call.
- Sizing budgets and forecasting rough ROAS ranges with stated assumptions.

The parts you keep

You keep the final call on budget, the risk tolerance, the client relationship, and the read on creative quality. AI can tell you the maths; it cannot feel the client's pressure or know that a 'loser' ad was actually killed by a bad audience. That context is your edge.

A daily buying workflow

Morning: paste yesterday's numbers, get a 5-line read plus three actions. Midday: use the hook machine to feed the creative team. Evening: draft tomorrow's test plan. This turns a buyer into a decision-maker with an assistant, not a spreadsheet-operator.

TIP

Never paste a client's confidential account name or identifiers into a free consumer tool. Anonymise the data first, or use an enterprise tier your company controls.

07 | MODULE

AI for Reporting and Analytics

Turn raw data into decisions your client or founder can act on, and stop spending your evenings building charts.

From export to insight in minutes

Paste your ad or analytics export and ask for: the three biggest changes week-over-week, the one insight that matters most, and one clear recommendation with the trade-off. The model does the summarising; you sanity-check the numbers.

Write the narrative, not just the numbers

- Turn a metrics table into a plain-English summary your client actually reads.
- Draft the weekly report structure: wins, losses, why, next steps.
- Anticipate the client's three likely questions and pre-answer them.
- Translate jargon: ROAS, CPM, MER, CTR into money language.

Forecasting with stated assumptions

Ask for a conservative, base, and aggressive forecast with the assumptions listed. The value is not the exact number; it is forcing the assumptions into the open so you and the client agree on what has to be true for the plan to work.

TIP

The analyst who can explain 'why' in one sentence is worth more than the one who can only produce the chart. AI lets you be the first while it does the chart.

08 | MODULE

AI for CRO and Landing Pages

Find the leaks in the funnel you are paying to fill, and fix them with copy and structure before you touch another rupee of spend.

Diagnose the page before you redesign it

Feed the model the landing page copy, the offer, and the traffic source, and ask for a CRO teardown: the clarity gaps, the trust gaps, and the friction points, ranked by likely conversion impact. You get a hypothesis list in minutes instead of a week.

Rewrite for conversion

- A headline that states the outcome, not the product.
- Benefit-first bullets that answer 'what do I get?'
- Objection handling drawn from real customer reviews.
- A clearer call to action and a risk-reversal offer.

Test like a marketer, not a guesser

Turn each hypothesis into an A/B test with a clear prediction: 'changing the headline to state the outcome will lift conversion because the current one is unclear.' That is the discipline that separates CRO from random editing.

TIP

A landing page fix that lifts conversion 20 percent is worth more than a 20 percent bigger ad budget, and it costs nothing to test.

Now go run it.

AI will not replace the performance marketer. A performance marketer who uses AI well will replace the one who does not. You just became the first kind.

Questions or feedback: adydaddy81@gmail.com

AdyDaddy | www.adydaddy.com | Rs 1,499 one-time